

What is a Database?

Cheat sheets, Practice Exams and Flash cards 🖱️ www.exampor.co/clf-c01

A database is a **data-store that stores semi-structured and structured data.**

A database is more **complex data stores** because it **requires using formal design and modeling techniques**

Databases can be generally categorized as either:

- **Relational databases**
 - Structured data that strongly represents tabular data (tables, rows and columns)
 - Row-oriented or Columnar-oriented
- **Non-relational databases**
 - Semi-structured that may or may not distantly resemble tabular data.

Databases have a rich set of functionality:

- specialized language to query (retrieve data)
- specialized modeling strategies to optimize retrieval for different use cases
- more fine tune control over the transformation of the data into useful data structures or reports



Normally a databases infers someone is using a **relational row-oriented data store**

What is Data Warehouse?

Cheat sheets, Practice Exams and Flash cards www.exampopro.co/clf-c01

A relational datastore designed for **analytic workloads**, which is generally **column-oriented data-store**

Companies will have **terabytes and millions of rows of data**, and they need a fast way to be able to produce analytics reports

Data warehouses generally perform **aggregation**


- aggregation is grouping data eg. find a total or average
- Data warehouses are optimized around columns since they need to quickly aggregate column data

Data warehouses are generally designed be HOT

- Hot means they can return queries very fast even though they have vast amounts of data

Data warehouses are infrequently accessed meaning they aren't intended for real-time reporting but maybe once or twice a a day or once a week to generate business and user reports.

A data warehouse needs to consume data from a relational databases on a regular basis.



What is a Key / Value store?

Cheat sheets, Practice Exams and Flash cards  www.exampopro.co/clf-c01

A **key-value database** is a type of non-relational database (NoSQL) that uses a simple key-value method to store data.

A key/value stores a **unique key** alongside a value



Key	Value
Data	1010101000101011001010010101001
Worf	0110101100010101010101011100010
Ro Laren	0010101001010110010101010101010

Key values stores are **dumb and fast**.

They generally lack features like:

- Relationships
- Indexes
- Aggregation

Key	Value
Data	{species: android, rank: 'Lt commander' }
Worf	{species: klingon, rank: 'Lt commander' }
Ro Laren	{species: bajoran, affiliation: 'maquis'}



A simple key/value store will interpret this data resembling a dictionary (aka Associative arrays or hash)

A key/value store can resemble tabular data, it does not have to have the consistent columns per row (hence its schemaless)



Key (Name)	Species	Rank	Affiliation
Data	andriod	Lt commander	
Worf	klingon	Lt commander	
Ro Laren	bajoran		maquis

Due to their simple design they can scale well beyond a relational database

What is a Document store?

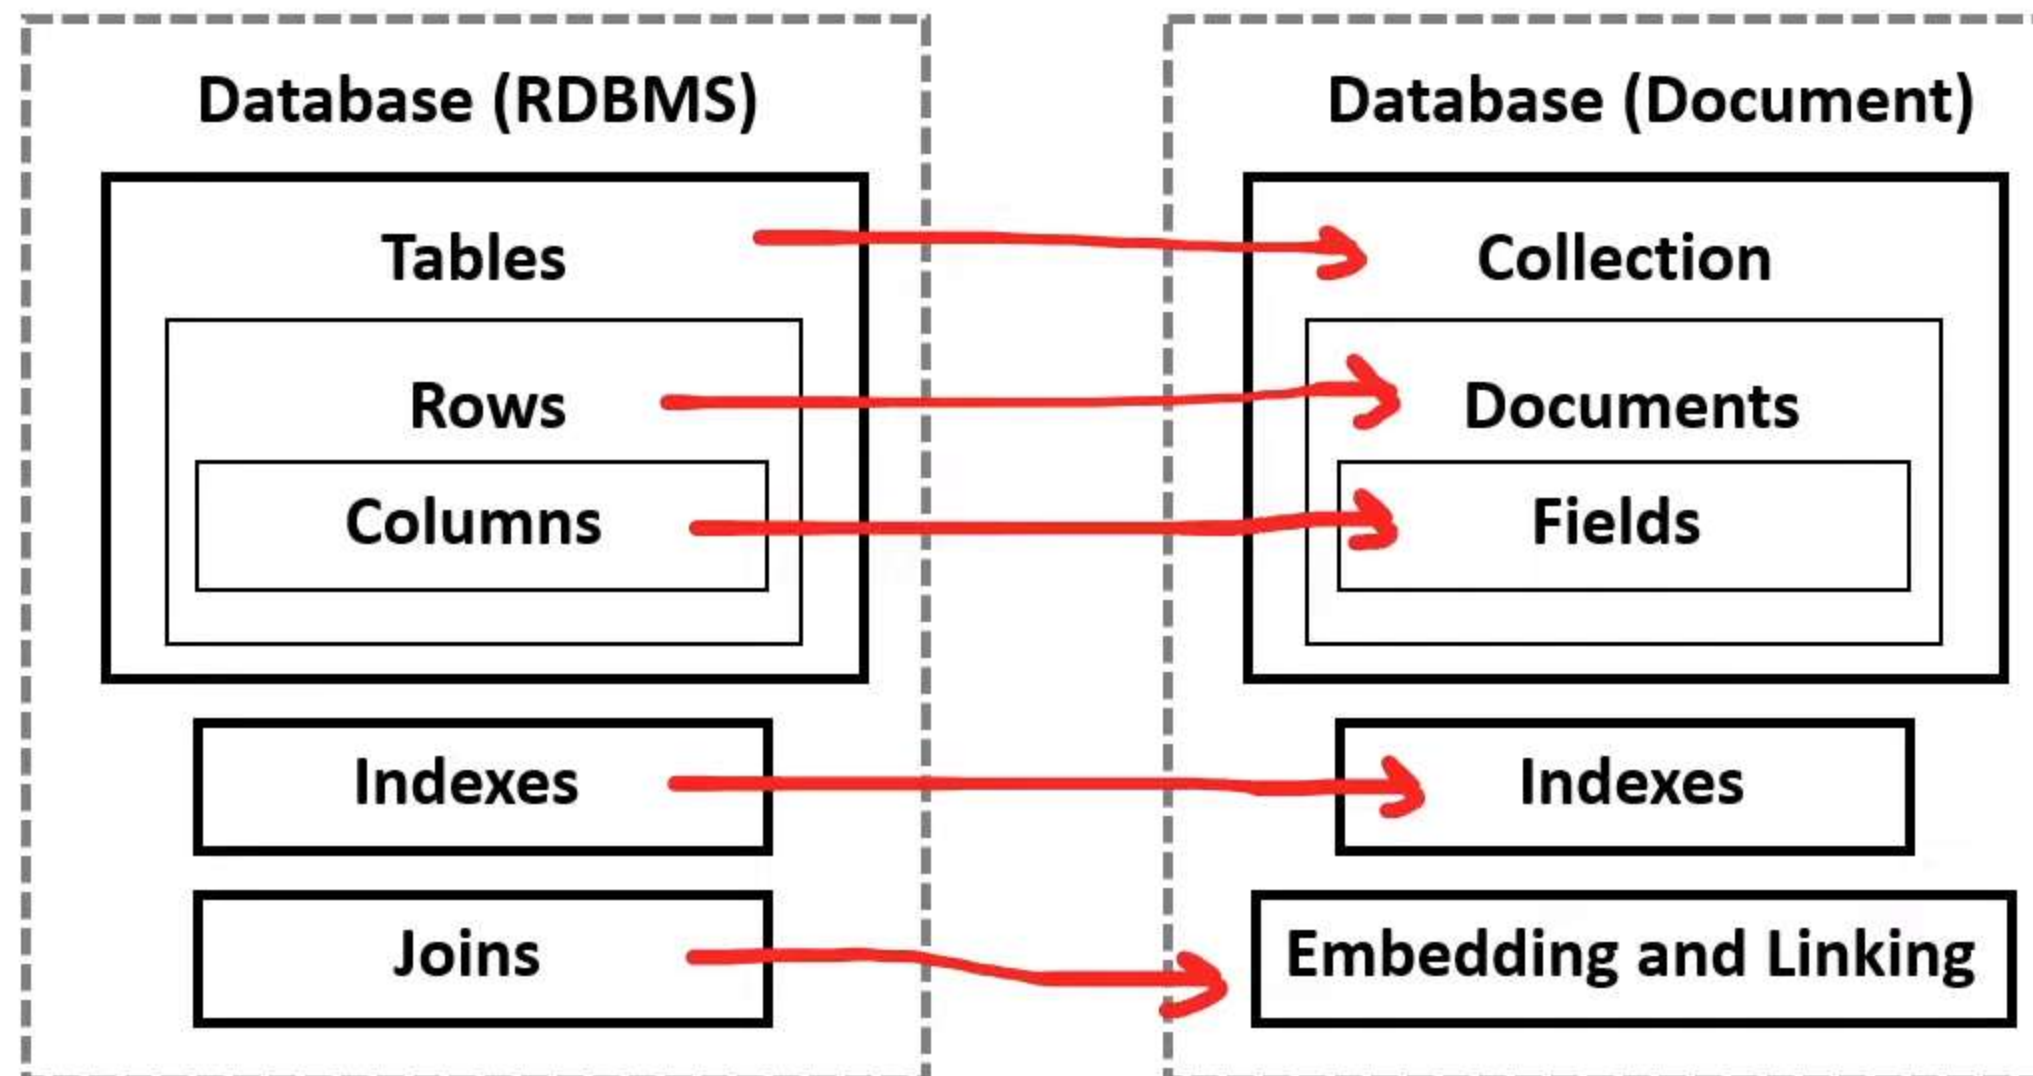
Cheat sheets, Practice Exams and Flash cards 🖱️ www.examprompro.co/clf-c01

A document store is a NOSQL database that stores **documents** as its primary data structure.

A document could be an XML but more commonly is JSON or JSON-Like

Document stores are sub-class of Key/Value stores

The components of a document store compared to Relational database



NoSQL Database Service

Cheat sheets, Practice Exams and Flash cards  www.exampor.co/clf-c01



DynamoDB is a serverless **NoSQL key/value and document database**. It is designed to scale to **billions of records** with guaranteed consistent data return in at least a second. You don't have to worry about managing shards!



DynamoDB is AWS's flagship database service meaning whenever we *think* of a database service that just scales, is cost effective and very fast we should think DynamoDB



In **2019**, **Amazon** the online shopping retail shutdown their last Oracle database and completed their migration to DynamoDB. They had 7,500 Oracle Database and 75 petabytes of data. With DynamoDB they reduce costs by 60% and reduce latency by 40%

When we want a massively scalable database



DocumentDB is a NoSQL **document** database that is "MongoDB compatible"

MongoDB is very popular NoSQL among developers. There were open-source licensing issues around using open-source MongoDB, so AWS got around it by just building their own MongoDB database.

When you want a MongoDB database.



Amazon Keyspaces is a fully managed Apache Cassandra database. Cassandra is an open-source NoSQL key/value database similar to DynamoDB in that is columnar store database but has some additional functionality. *When you want to use Apache Casandra.*



Relational Database Services

Cheat sheets, Practice Exams and Flash cards  www.examprompro.co/clf-c01



Relational Database Service (RDS) is a **relational database service** that supports multiple SQL engines. Relational is synonymous with SQL and Online Transactional Processing (OLTP). Relational database are **the most commonly used type of database** among tech companies and start-ups.

RDS Supports the following SQL Engines:

- **MySQL** – The most popular open-source SQL database that was purchased and now owned by Oracle.
- **MariaDB** – When Oracle bought MySQL. MariaDB made a fork (copy) of MySQL was made under a different open-source license.
- **Postgres (PSQL)** – Most popular open-source SQL database among developers. Has rich-features over MySQL but at added complexity
- **Oracle** – Oracle's proprietary SQL database. Well used by Enterprise companies. You have to buy a license to use it.
- **Microsoft SQL Server** – Microsoft's proprietary SQL database. You have to buy a license to use it.
- **Aurora** – Fully managed database.



Aurora is a **fully managed** database of either MySQL (5x faster) and PSQL (3x faster) database.

When you want a highly available, durable, scalable and secure relational database for Postgres or MySQL



Aurora Serverless is the **serverless on-demand version of Aurora**. *When you want "most" of the benefits of Aurora but can trade to have cold-starts or you don't have lots of traffic demand*



RDS on VMware allows you to deploy RDS supported engines to on an-premise data-center. The datacenter must be using VMware for server virtualization. *When you want databases managed by RDS on your own datacenter*



Other Database Services

Cheat sheets, Practice Exams and Flash cards 🖱️ www.exampor.co/clf-c01



Redshift is a **petabyte-size data-warehouse**. Data-warehouses are for Online Analytical Processing (OLAP) Data-warehouses can be expensive because they are keeping data “hot”. Meaning that we can run a very complex query and a large amount of data and get that data back very fast.

When you to quickly generate analytics or reports from a large amount of data.



ElastiCache is a managed database of the **in-memory** and **caching** open-source databases **Redis** or **Memcached**. *When you need to improve the performance of application by adding a caching layer in-front of web-server or database.*



Neptune is a managed graph database. Data is represented as interconnected nodes.

When you need to understand the connections between data eg. Mapping Fraud Rings or Social Media relationships



Amazon Timestreams is a fully managed time series database. Think of devices that send lots of data that are time-sensitive such as IoT devices. *When you need to measure how things change over time.*



Amazon Quantum Ledger Database is a fully managed ledger database that provides transparent, immutable and cryptographically variable transaction logs.

When you need to record history of financial activities that can be trusted.



Database Migration Service (DMS) is database migration service. You can migrate from:

- on-premise database to AWS
- from two database in different or same AWS accounts using different SQL engines
- from an SQL to NoSQL database